
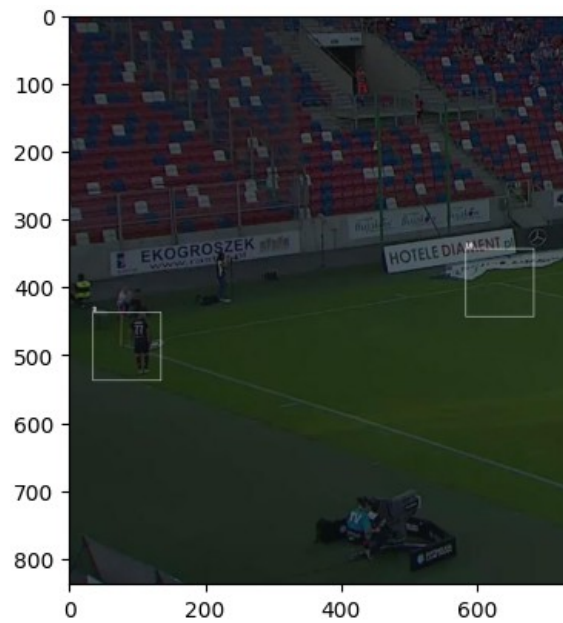
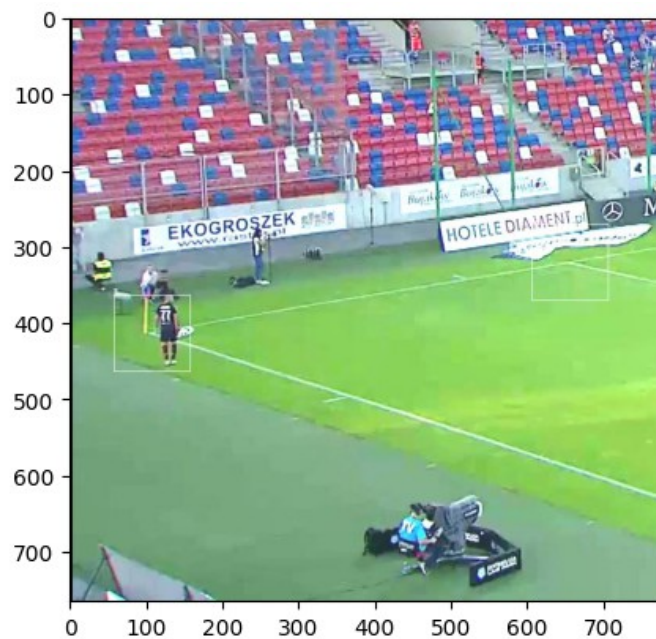
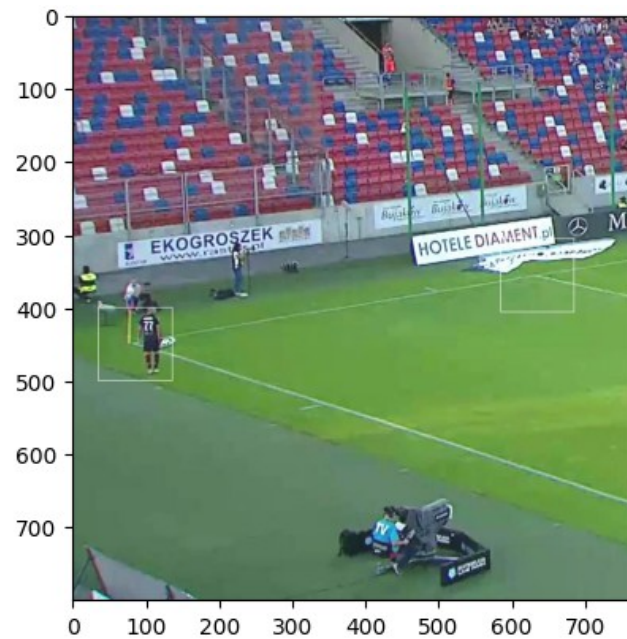


Experiment 1 – Data Augmentation

Data augmentation pipeline is included in Jupyter Notebook titled '4_Data_Augmentation.ipynb'. To increase samples of classes 3 and 23 both during training and validation (and testing in this case), clipping and brightening of images was applied.

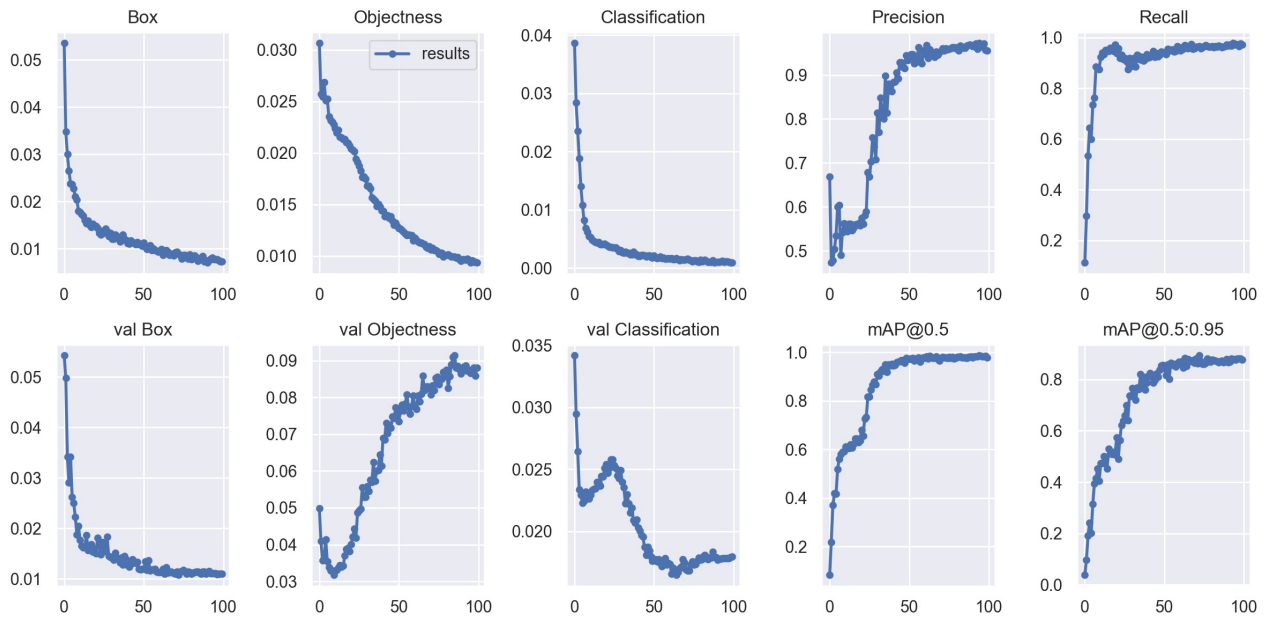
- For each image containing classes 3 and 10 or 23 and 30 respectively, 4 images with increased or decreased brightness and different x,y of labelled class is generated – like below:








- 72 and 16 additional images containing classes 3 and 10 were respectively added to training set and validation set.
- 108 and 28 additional images containing classes 23 and 30 were respectively added to training set and validation set.


Unfortunately it did not improve performance of the model, although considering shapes of training charts (val Objectness – it started to fall near epoch 100 and classification – it is possible it starts to fall in next few epochs), new model might be underfitted. I think that this experiment needs another 100 epochs.



Epoch	gpu_mem	box	obj	cls	total	labels	img_size	
99/99	9.49G	0.007318	0.009371	0.0009104	0.0176	92	640: 100%	
	Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95: 100%	
	all	796	6696	0.956	0.972	0.979	0.876	
	1	796	218	0.988	0.972	0.982	0.931	
	2	796	368	0.997	0.997	0.996	0.948	
	3	796	20	0.872	0.685	0.843	0.684	
	4	796	109	0.967	0.991	0.982	0.911	
	5	796	202	0.986	0.975	0.982	0.898	
	6	796	266	0.978	0.977	0.992	0.945	
	7	796	246	0.977	0.98	0.992	0.929	
	8	796	203	0.984	0.99	0.993	0.913	
	9	796	127	0.961	0.982	0.993	0.931	
	10	796	47	0.958	0.977	0.976	0.804	
	11	796	164	0.987	0.976	0.976	0.886	
	12	796	193	1	0.984	0.995	0.908	
	13	796	112	0.982	0.972	0.994	0.922	
	14	796	80	0.987	0.943	0.99	0.93	
	15	796	187	0.949	0.973	0.981	0.909	
	16	796	380	0.996	0.997	0.994	0.937	
	17	796	382	0.996	0.997	0.998	0.934	
	18	796	372	0.993	0.997	0.995	0.956	
	19	796	251	0.971	0.976	0.99	0.603	
	20	796	372	0.965	0.995	0.995	0.866	
	21	796	363	0.989	0.981	0.997	0.866	
	22	796	184	0.962	0.989	0.985	0.935	
	23	796	35	0.784	0.971	0.854	0.668	
	25	796	164	0.963	0.982	0.991	0.904	
	26	796	232	0.973	0.974	0.994	0.942	
	27	796	215	0.977	0.985	0.994	0.937	
	28	796	175	0.946	0.994	0.997	0.904	
	29	796	112	0.929	1	0.993	0.95	
	30	796	48	0.8	0.896	0.889	0.652	
	31	796	145	0.986	1	0.996	0.899	
	32	796	166	0.971	1	0.996	0.919	
	33	796	106	0.951	0.981	0.993	0.906	
	34	796	72	0.857	0.986	0.98	0.901	
	35	796	161	0.924	0.988	0.982	0.919	
	39	796	219	0.954	0.958	0.987	0.616	


epochs completed in 11.250 hours.






Experiment 2 – Bigger bounding box for some classes

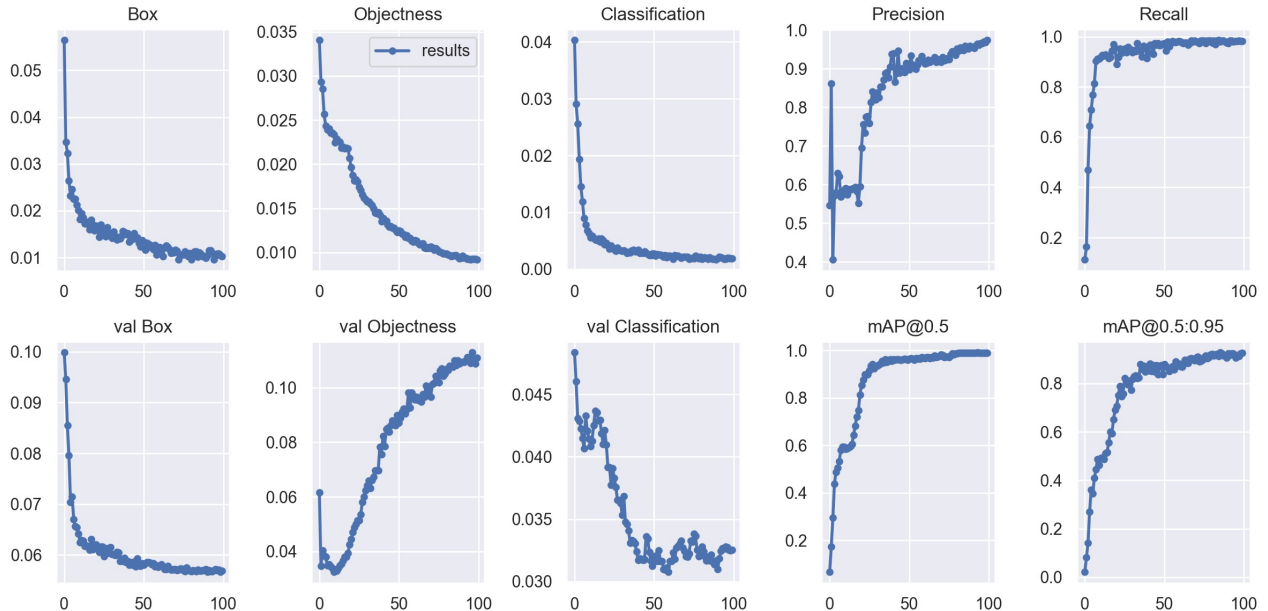
Data augmentation pipeline is included in Jupyter Notebook titled '4_Data_Augmentation.ipynb'. This experiment featured bigger bounding boxes (by 60 %) for classes: 3, 23, 10, 30, 39, 19, 20, 21.



9677b5248c7d2a391bb331263538c8







This model might be underfitted too and perhaps more epochs would help, but the results are a bit better and definitely better balanced (based on $mAP@.5:.95$). On the other hand average precision for classification is slightly lower – I think it is underfitted especially since precision is still increasing.



Epoch	gpu_mem	box	obj	cls	total	labels	img_size	
99/99	7.44G	0.01024	0.009215	0.00185	0.0213	12	640:	100%
Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95:		
all	765	6608	0.974	0.982	0.989	0.925		
1	765	218	0.958	0.968	0.989	0.96		
2	765	368	0.981	0.997	0.996	0.983		
3	765	4	0.999	1	0.995	0.995		
4	765	109	0.954	0.991	0.992	0.96		
5	765	202	0.99	0.968	0.977	0.924		
6	765	266	0.99	0.974	0.994	0.962		
7	765	246	0.987	0.988	0.989	0.944		
8	765	203	0.995	0.983	0.99	0.924		
9	765	127	0.992	0.972	0.995	0.961		
10	765	31	0.954	0.968	0.982	0.96		
11	765	164	0.992	0.957	0.983	0.926		
12	765	193	0.994	0.984	0.995	0.946		
13	765	112	0.991	0.963	0.995	0.95		
14	765	80	0.987	0.979	0.991	0.958		
15	765	187	0.983	0.95	0.978	0.915		
16	765	380	0.994	0.997	0.995	0.974		
17	765	382	0.997	0.997	0.997	0.959		
18	765	372	0.995	0.997	0.995	0.984		
19	765	251	0.992	0.985	0.996	0.744		
20	765	372	0.955	0.981	0.993	0.915		
21	765	363	0.978	0.964	0.995	0.914		
22	765	184	0.973	0.978	0.985	0.956		
23	765	7	0.958	1	0.995	0.936		
25	765	164	0.97	0.971	0.987	0.925		
26	765	232	0.982	0.983	0.994	0.963		
27	765	215	0.98	0.991	0.995	0.956		
28	765	175	0.994	0.989	0.996	0.927		
29	765	112	0.956	1	0.992	0.951		
30	765	20	0.867	0.98	0.943	0.681		
31	765	145	0.96	0.985	0.995	0.919		
32	765	166	0.986	1	0.99	0.929		
33	765	106	0.928	0.991	0.99	0.905		
34	765	72	0.945	0.986	0.978	0.915		
35	765	161	0.963	0.966	0.983	0.927		
39	765	219	0.964	0.979	0.988	0.733		

100 epochs completed in 11.602 hours.





After testing with my metrics we can see that the results are much worse for particular classes. Model is probably underfitted. Note that if any problem with precision of bounday box attachment occurs, MSE increases drastically.

